

STA 322 Project 1: Characteristics of Trinity College Faculty

Janet Meng and Amy Duan

2026-10-09

Introduction

This project uses a probability sample to estimate characteristics of faculty in Duke's Trinity College of Arts & Sciences. We answer four questions, each with a point estimate and a 95% confidence interval:

1. What fraction of faculty are women?
2. What fraction of faculty have at least one grant that is ongoing in 2026?
3. Among faculty with a Ph.D., what is the average number of years since the degree?
4. Do women and men differ in the fraction who are full professors?

The sampling frame is built from departmental web pages, and the data for each sampled faculty member come from their Scholars@Duke profile. Although it might be possible to collect information on every faculty member, the purpose of the project is to practice survey sampling, so we drew a probability sample and use design-based estimation.

Target Population and Sampling Frame

The target population is everyone listed on the faculty web pages of Trinity College's 36 departments and programs as of October 1, 2026. For each unit we used the "primary faculty" page where one exists, because it lists faculty whose main appointment is in that department. Some units do not post a primary-faculty page. For those (e.g., English, Theater Studies, Music, Religious Studies), we used the full faculty listing. The population therefore includes tenure-track faculty, lecturers, instructors, research and practice faculty, ROTC cadre, and any emeriti whom a department still lists.

Because some departments list all of their faculty rather than only primary faculty, a faculty member with a joint appointment can appear on more than one roster. Removing these duplicates would require matching every name across all 36 rosters before sampling. We therefore retained every listing and treated each faculty member appearing on a department's page as a member of that department. The practical consequence is that a person listed on two rosters has two chances of selection and could, in principle, be drawn twice, once through each listing. We flag faculty who appear on more than one roster in our data and discuss this as a limitation.

We built the frame by counting each roster by hand and recording, for each unit, the number of faculty listed. This gives a frame of $N = 866$ faculty across 36 units.

Survey Design

Overview

We used **stratified simple random sampling without replacement with proportional allocation**, with a total sample size of $n = 100$. The population is partitioned into $H = 27$ strata: departments, with a few very small units pooled. Within each stratum h we selected a simple random sample of n_h roster positions out of the N_h faculty listed.

The sample-size calculation, the full allocation table, and the derivation of the weights and variance estimator are given in the Appendix.

Defining the strata

Each department is its own stratum unless it is too small to receive at least two sampled faculty under proportional allocation. With proportional allocation, stratum h receives nN_h/N units before rounding, so a department needs

$$\frac{nN_h}{N} \geq 2 \iff N_h \geq \frac{2N}{n} = \frac{2 \times 866}{100} \approx 17.3$$

faculty to get two sampled people on its own. Units below this size were pooled with related units:

- **Classical & Religious Studies:** Classical Studies, Religious Studies
- **Dance & Theater:** Dance, Theater Studies
- **Education & Gender Studies:** Education; Gender, Sexuality & Feminist Studies
- **Small Languages:** German Studies, Romance Studies, Slavic & Eurasian Studies, International Comparative Studies
- **Military & PE:** Air Force ROTC, Army ROTC, Navy ROTC, Health, Wellness & Physical Education

We pooled related fields so that each pooled stratum is still fairly homogeneous, which preserves most of the benefit of stratification.

Implementing the design in R

The `survey` package needs four pieces of information about each sampled faculty member to reproduce this design:

- **ids = ~1:** there is no clustering; each faculty member is selected individually.
- **strata = ~stratum:** the stratum h the person was sampled from, so variances are computed within strata and then combined.
- **weights = ~weight:** the design weight $1/\pi_i = N_h/n_h$.
- **fpc = ~N_h:** the stratum population size N_h , from which `survey` computes n_h/N_h and applies the finite population correction $1 - n_h/N_h$.

```
des <- svydesign(ids = ~1, strata = ~stratum, weights = ~weight,
               fpc = ~N_h, data = dat)
```

The estimator and variance formula that `survey` computes from these inputs are derived in Appendix C, and the full code is in Appendix D.

Rationale for Our Design Choices

This section explains why we chose to stratify by department and why we allocated the sample proportionally.

Why stratify by department

We stratified by department mainly for precision. The share of women and the prevalence of grants differ sharply across fields, and when the stratum means $\bar{Y}_h$ are spread out, the within-stratum variances S_h^2 are small relative to S^2 , so stratified sampling has smaller variance than a simple random

sample of the same size. Stratifying also guarantees that every department appears in the sample, so no field can be missed by chance. Finally, it was convenient: the frame is already organized by department, so the stratum sizes N_h came directly from counting each roster, and sampling within a stratum only requires that department’s page.

Why proportional allocation

We allocated the sample proportionally, $n_h \approx nN_h/N$, rounding so the total is exactly 100. Proportional allocation guarantees $\text{Var}(\hat{t}_{prop}) \leq \text{Var}(\hat{t}_{srs})$, because it removes the between-strata variation from the variance. Unlike optimal allocation, which needs the unknown S_h and is tuned to a single variable, it requires only the stratum sizes N_h and serves all four questions equally. It is also nearly self-weighting: every faculty member’s selection probability is $\pi_i = n_h/N_h \approx n/N$.

The resulting n_h , inclusion probabilities, and weights for all 27 strata are listed in Appendix B.

Data Collection

Drawing the sample

We drew the sample in R with `set.seed(322)` so it can be reproduced exactly. In each stratum h we selected n_h distinct roster positions at random from $1, \dots, N_h$. In pooled strata, positions run through the member departments in a fixed order, and each position was mapped back to a department and a position on that department’s page. The sampled person is the name at that position, counting down the page in the order shown on the website. The sampling code is reproduced in Appendix D. After drawing the sample, we checked the sampled names for anyone selected more than once; there were none.

Variables and coding rules

For each sampled person we recorded the following from their Scholars@Duke profile:

- **Sex (Question 1).** Coded from pronouns in the profile biography where given, and otherwise from first name and profile photo.
- **Ongoing grant in 2026 (Question 2).** Coded 1 if any grant listed on the profile has an award period that overlaps calendar year 2026, regardless of the person’s role (principal investigator, co-investigator, or mentor on a training grant). Fellowships and gifts, which Scholars lists separately from grants, were not counted.
- **Ph.D. status and year (Question 3).** From the Education section. A D.Phil. was counted as a Ph.D.. An M.F.A., M.Ed., M.A., or other degree was not. Years since Ph.D. is $2026 - \text{Ph.D. year}$.
- **Rank (Question 4).** We recorded the full title and coded *full professor* as 1 for any title at the rank of Professor, including named and distinguished chairs, Professor of the Practice, Research Professor, and Emeriti at that rank. Assistant and associate professors, lecturers, instructors, and other titles were coded 0.

Missing data and frame issues

- **Missing information.** A few profiles were incomplete or missing information, which was coded as missing values. Missing values are excluded from the estimate for that question only and were not imputed because the small number of missing values has no practical effect on any estimate. No values of sex, Ph.D. status, or rank are missing, and Ph.D. year is blank only for the 14 faculty without a Ph.D..

- **Emeriti.** Of the sampled faculty, 6 are emeriti who are still listed on a department roster. They belong to our target population as defined, so we kept them rather than replacing them, since replacing sampled units would change the selection probabilities.
- **Roster changes.** The Gender, Sexuality & Feminist Studies page listed 8 faculty when we built the frame and 9 when we collected data. We used the person at the sampled position on the current page and treat this as a minor frame error.
- **Multiple listings.** Of the sampled faculty, 11 were flagged as also appearing on another department's roster (for example, faculty with joint appointments in African & African American Studies or Theater Studies). We checked the sampled names, and no faculty member was selected more than once, so each of the 100 records is a distinct person. Anyone repeating this procedure should run the same check and report any repeats, because any characteristic associated with joint appointments may be slightly overrepresented in their estimates.

Before analysis we checked that every record had a valid stratum, that coded variables took only the values 0 and 1, and that a Ph.D. year was present only for faculty coded as having a Ph.D. The final data set has 100 faculty, 86 with a Ph.D. and 14 without.

Results and Analysis

All estimates below use the design weights and the stratified Horvitz-Thompson estimator with the finite population correction (Appendix C). Proportions are estimated with `svyciprop`, and the mean in Question 3 with `svymean` on the subset of Ph.D. holders.

Table 1: Design-based estimates with 95% confidence intervals ($n = 100$).

Question	Estimate	95% CI
1. Women	44.4%	36.2% to 52.9%
2. Ongoing grant in 2026	29.1%	21.5% to 38.2%
3. Mean years since Ph.D. (Ph.D. holders)	23.6 years	20.8 to 26.4 years
4a. Full professor, all faculty	52.3%	42.7% to 61.6%
4b. Full professor, women	52.2%	37.3% to 67.2%
4c. Full professor, men	52.3%	40.2% to 64.4%
4d. Difference in full professors (women - men)	-0.1%	-19.4% to 19.3%

Question 1: Fraction of faculty who are women

We estimate that 44.4% of Trinity faculty are women (95% CI: 36.2% to 52.9%). The sample contained 44 women and 56 men. The interval includes 50%, so the data do not rule out an even split, but they are more consistent with women being a minority of the faculty.

Question 2: Ongoing grants in 2026

We estimate that 29.1% of faculty have at least one grant ongoing in 2026 (95% CI: 21.5% to 38.2%). Most faculty therefore do not have an active grant. We counted a grant if its award period overlapped 2026, regardless of the faculty member's role, so training grants where the person is a mentor are included. A stricter definition (principal investigator only) would give a lower fraction.

Question 3: Years since Ph.D.

Among Ph.D. holders, the average time since the degree is 23.6 years (95% CI: 20.8 to 26.4). We estimate that about 86% of faculty hold a Ph.D. (95% CI: 79.6% to 90.0%), and the sample's Ph.D. years range from 0 to 57 years ago. The range is large because the faculty include recent hires and senior professors.

Question 4: Do women and men differ in rank?

Our “full professor” code groups ranks that differ in meaning (distinguished and endowed chairs, Professors of the Practice, and emeriti with the rank). We estimate that 52.3% of faculty are full professors overall (95% CI: 42.7% to 61.6%). For women the estimate is 52.2% and for men 52.3%. The estimated difference is -0.1% (95% CI: -19.4 to 19.3), and the Rao-Scott test gives $p = 0.99$. We therefore find no evidence of a difference in the share of full professors between women and men. This comparison does not adjust for years since the Ph.D., so any real gap in rank may be hidden or explained by seniority.




Figure 1: Estimated proportions with 95% confidence intervals.

Limitations

- **Measurement of sex.** Sex was coded using pronouns in the biography where available and otherwise first name and photo. A few profiles state no sex identifiers, so some codes are inferences that may not match how those faculty identify. Misclassification would bias Question 1 slightly and would also affect the Question 4 comparison.
- **Frame and duplicate listings.** Faculty with joint appointments can appear on more than one roster. Those faculty would have more than one chance of selection but receive the weight of the single department they were drawn from, so estimates could be slightly biased toward their characteristics. That did not happen here, but it could in a repeated draw, so a replication should check for repeated names. A cleaner fix would be to deduplicate the frame by assigning each person to one primary department before sampling, or to divide each person's weight by their number of listings.
- **Design.** Several strata have only two or three sampled faculty, so the stratum-level variance estimates are unstable even though the overall variance is sound. Pooling unrelated small departments, though necessary, gives those strata less meaningful interpretations.

- **Missing data.** The one missing grant value was excluded, not imputed. This has negligible effect here but assumes that faculty who disable their profile are not systematically different.

Conclusion

We used a stratified simple random sample of 100 faculty, allocated in proportion to department size, to estimate characteristics of Trinity College faculty. We estimate that about 44% of faculty are women (95% CI: 36.2% to 52.9%), that about 29% have a grant ongoing in 2026 (95% CI: 21.5% to 38.2%), and that Ph.D. holders earned their degree an average of 24 years ago (95% CI: 20.8 to 26.4). For our additional question, the share of full professors is nearly identical for women (52%) and men (52%), with no statistically detectable difference, although the sample is too small to rule out a moderate one.

The main limits are the width of the confidence intervals at $n = 100$, the hand-coded sex and grant variables, and the potential for duplicate roster listings (although there are none in our sample). With more time we would increase the sample, run a pilot to estimate department-level variability for optimal allocation, and measure rank more finely and adjust for years since the Ph.D. to learn whether women and men differ in rank once seniority is accounted for.

Appendix

A. Sample size

We chose $n = 100$, comfortably above the suggested 75 and manageable to collect in a few hours each. We sample $n/N = 100/866 \approx 0.115$ of the population. Ignoring the gain from stratification, the worst-case margin of error for a proportion (when the proportion is 0.5, so $S^2 \approx 0.25$) is

$$1.96\sqrt{\left(1 - \frac{n}{N}\right) \frac{S^2}{n}} = 1.96\sqrt{\left(1 - \frac{100}{866}\right) \frac{0.25}{100}} \approx 0.092,$$

about ± 9 percentage points. Stratification should make the actual intervals somewhat narrower.

B. Allocation and weights

Because we take a simple random sample within each stratum, faculty member i in stratum h has inclusion probability

$$\pi_i = \Pr(I_i = 1) = \frac{n_h}{N_h},$$

and their **design weight** is the inverse, $1/\pi_i = N_h/n_h$. The weight is the number of population faculty that each sampled person represents. The weights within a stratum sum to N_h , so across all strata $\sum_{i \in S} 1/\pi_i = N = 866$.

Under proportional allocation, $n_h \approx nN_h/N$. Because nN_h/N is rarely a whole number, we used largest-remainder rounding: round every n_h down, then give the remaining units, one each, to the strata with the largest fractional parts, so the total is exactly 100. The code below computes the strata, the allocation, and the weights from our roster counts.

```
frame <- read.csv("faculty_counts.csv", stringsAsFactors = FALSE)
N <- sum(frame$total_faculty)
n <- 100

# Pool units too small to receive 2 sampled faculty on their own
pool <- list(
  "Military & PE (pooled)" = c("Air Force ROTC", "Army ROTC", "Navy ROTC",
    "Health, Wellness & Physical Education"),
  "Small languages (pooled)" = c("International Comparative Studies",
    "Slavic & Eurasian Studies", "Romance Studies",
    "German Studies"),
  "Dance & Theater (pooled)" = c("Dance", "Theater Studies"),
  "Education & Gender Studies (pooled)" = c("Education",
    "Gender, Sexuality & Feminist Studies"),
  "Classical & Religious Studies (pooled)" = c("Classical Studies",
    "Religious Studies")
)
frame$stratum <- frame$department
for (s in names(pool)) frame$stratum[frame$department %in% pool[[s]]] <- s

strata <- aggregate(total_faculty ~ stratum, data = frame, FUN = sum)
names(strata)[2] <- "N_h"
```

```

# Proportional allocation with largest-remainder rounding
strata$exact <- n * strata$N_h / N
strata$n_h <- floor(strata$exact)
short <- n - sum(strata$n_h)
bump <- order(strata$exact - strata$n_h, decreasing = TRUE)[seq_len(short)]
strata$n_h[bump] <- strata$n_h[bump] + 1

# Inclusion probabilities and design weights
strata$pi <- strata$n_h / strata$N_h
strata$weight <- strata$N_h / strata$n_h

stopifnot(sum(strata$n_h) == n, all(strata$n_h >= 2),
           isTRUE(all.equal(sum(strata$n_h * strata$weight), N)))

```

Table 2: Proportional allocation and design weights by stratum

Stratum	N_h	nN_h/N	n_h	$\pi_i = n_h/N_h$	$1/\pi_i = N_h/n_h$
African & African American Studies	18	2.08	2	0.111	9.00
Art, Art History & Visual Studies	27	3.12	3	0.111	9.00
Asian & Middle Eastern Studies	45	5.20	5	0.111	9.00
Biology	61	7.04	7	0.115	8.71
Chemistry	27	3.12	3	0.111	9.00
Classical & Religious Studies (pooled)	27	3.12	3	0.111	9.00
Computer Science	38	4.39	4	0.105	9.50
Cultural Anthropology	17	1.96	2	0.118	8.50
Dance & Theater (pooled)	28	3.23	3	0.107	9.33
Economics	60	6.93	7	0.117	8.57
Education & Gender Studies (pooled)	20	2.31	2	0.100	10.00
English	33	3.81	4	0.121	8.25
Evolutionary Anthropology	17	1.96	2	0.118	8.50
History	27	3.12	3	0.111	9.00
Linguistics	20	2.31	2	0.100	10.00
Literature	17	1.96	2	0.118	8.50
Mathematics	58	6.70	7	0.121	8.29
Military & PE (pooled)	23	2.66	3	0.130	7.67
Music	65	7.51	7	0.108	9.29
Philosophy	18	2.08	2	0.111	9.00
Physics	40	4.62	5	0.125	8.00
Political Science	32	3.70	4	0.125	8.00
Psychology & Neuroscience	41	4.73	5	0.122	8.20
Small languages (pooled)	25	2.89	3	0.120	8.33
Sociology	25	2.89	3	0.120	8.33

Stratum	N_h	nN_h/N	n_h	$\pi_i = n_h/N_h$	$1/\pi_i = N_h/n_h$
Statistical Science	26	3.00	3	0.115	8.67
Thompson Writing Program	31	3.58	4	0.129	7.75

The 27 strata have n_h between 2 and 7, and the weights range from 7.67 to 10, close to the overall $N/n = 8.66$. Rounding is the only source of the differences: strata rounded down (e.g., Education & Gender Studies, $2.31 \rightarrow 2$) have slightly larger weights, and strata rounded up (e.g., Military & PE, $2.66 \rightarrow 3$) have slightly smaller ones.

C. Estimation and variance

Plugging the inclusion probabilities into the Horvitz-Thompson estimator gives the stratified estimators of the population total and mean:

$$\hat{t}_y = \sum_{i \in S} \frac{y_{hi}}{\pi_i} = \sum_{h=1}^H N_h \bar{y}_h, \quad \bar{y}_{HT} = \frac{\hat{t}_y}{N} = \sum_{h=1}^H \frac{N_h}{N} \bar{y}_h.$$

Each stratum's sample mean $\bar{y}_h$ is weighted by that stratum's share of the population. For a 0/1 variable such as sex, $\bar{y}_h$ is the sample proportion in stratum h , and $\bar{y}_{HT}$ estimates the population proportion. Because $\hat{t}_y$ is a Horvitz-Thompson estimator, it is unbiased.

Because we take separate simple random samples in each stratum, the $\bar{y}_h$ are independent, and the variance of $\bar{y}_{HT}$ is

$$\text{Var}(\bar{y}_{HT}) = \sum_{h=1}^H \frac{N_h^2}{N^2} \left(1 - \frac{n_h}{N_h}\right) \frac{S_h^2}{n_h}, \quad S_h^2 = \sum_{i=1}^{N_h} \frac{(y_{hi} - \bar{Y}_h)^2}{N_h - 1}.$$

We estimate each S_h^2 with the sample variance s_h^2 , giving

$$\widehat{\text{Var}}(\bar{y}_{HT}) = \sum_{h=1}^H \frac{N_h^2}{N^2} \left(1 - \frac{n_h}{N_h}\right) \frac{s_h^2}{n_h},$$

and the 95% confidence interval $\bar{y}_{HT} \pm 1.96 \sqrt{\widehat{\text{Var}}(\bar{y}_{HT})}$. For a 0/1 variable, $s_h^2 = n_h \bar{y}_h(1 - \bar{y}_h)/(n_h - 1)$.

Two parts of this formula shaped the design. The factor $1 - n_h/N_h$ is the **finite population correction**. We sample about 11.5% of each stratum without replacement, so this factor reduces the variance by about that much, and leaving it out would overstate our uncertainty. The divisor $n_h - 1$ in s_h^2 is why every stratum needs at least two sampled units.

Question 3 (average years since Ph.D. among Ph.D. holders) is a **domain estimate**: the number of Ph.D. holders in each stratum is not known in advance, so the estimate is a ratio of two estimated totals. We compute it by restricting the full design object to Ph.D. holders, which keeps the correct stratified variance.

D. Code

Drawing the sample

This code drew the roster positions used for data collection. It is shown, not re-run, so the report does not overwrite the sample file.

```

set.seed(322)
rows <- list()
for (i in seq_len(nrow(strata))) {
  s <- strata$stratum[i]
  members <- frame[frame$stratum == s, ]
  pos <- sort(sample(strata$N_h[i], strata$n_h[i]))
  ends <- cumsum(members$total_faculty)
  starts <- ends - members$total_faculty
  idx <- findInterval(pos - 1, ends) + 1
  rows[[i]] <- data.frame(stratum = s, department = members$department[idx],
                        roster_position = pos - starts[idx])
}
sample_ids <- do.call(rbind, rows)
sample_ids$weight <- strata$weight[match(sample_ids$stratum, strata$stratum)]
write.csv(sample_ids, "sample_ids.csv", row.names = FALSE)

```

Analysis

The code that builds the design object and computes every estimate in the Results and Analysis section:

```

library(survey)
knitr::opts_chunk$set(echo = FALSE, message = FALSE, warning = FALSE)

# Rebuild the strata exactly as in analysis.Rmd (pooled small units)
frame <- read.csv("faculty_counts.csv", stringsAsFactors = FALSE)
pool <- list(
  "Military & PE (pooled)" = c("Air Force ROTC", "Army ROTC", "Navy ROTC",
                              "Health, Wellness & Physical Education"),
  "Small languages (pooled)" = c("International Comparative Studies",
                              "Slavic & Eurasian Studies", "Romance Studies",
                              "German Studies"),
  "Dance & Theater (pooled)" = c("Dance", "Theater Studies"),
  "Education & Gender Studies (pooled)" = c("Education",
                              "Gender, Sexuality & Feminist Studies"),
  "Classical & Religious Studies (pooled)" = c("Classical Studies",
                              "Religious Studies"))

frame$stratum <- frame$department
for (s in names(pool)) frame$stratum[frame$department %in% pool[[s]]] <- s
N_h <- tapply(frame$total_faculty, frame$stratum, sum)

dat <- read.csv("data.csv", na.strings = c("NA", ""), stringsAsFactors = FALSE)
stopifnot(all(dat$stratum %in% names(N_h)))
n_h <- table(dat$stratum)
dat$N_h <- as.numeric(N_h[dat$stratum])
dat$weight <- dat$N_h / as.numeric(n_h[dat$stratum])
dat$years_since_phd <- ifelse(dat$has_phd %in% 1, 2026 - dat$phd_year, NA)

```

```

dat$sex <- factor(dat$female, levels = c(0, 1), labels = c("Men", "Women"))
dat$emeritus <- grepl("meritus", dat$rank_title) | grepl("^Emeritus", dat$notes)
dat$double_listed <- grepl("roster", dat$notes, ignore.case = TRUE) & !dat$emeritus
des <- svydesign(ids = ~1, strata = ~stratum, weights = ~weight,
               fpc = ~N_h, data = dat)

# helpers for inline numbers
pct <- function(x, d = 1) sprintf("%.*f%%", d, 100 * as.numeric(x))
ci_pct <- function(ci) sprintf("%.1f%% to %.1f%%", 100 * ci[1], 100 * ci[2])

# estimates
d1 <- subset(des, !is.na(female))
q1 <- svyciprop(~female, d1); q1_ci <- confint(q1)
d2 <- subset(des, !is.na(ongoing_grant_2026))
q2 <- svyciprop(~ongoing_grant_2026, d2); q2_ci <- confint(q2)
d3 <- subset(des, has_phd %in% 1 & !is.na(years_since_phd))
q3 <- svymean(~years_since_phd, d3); q3_ci <- confint(q3)
q3b <- svyciprop(~I(has_phd == 1), subset(des, !is.na(has_phd)))
q3b_ci <- confint(q3b)
d4 <- subset(des, !is.na(female) & !is.na(full_professor))
q4 <- svyciprop(~full_professor, d4); q4_ci <- confint(q4)
q4_by <- svyby(~full_professor, ~sex, d4, svymean, vartype = c("se", "ci"),
              covmat = TRUE)
q4_diff <- svycontrast(q4_by, c(-1, 1))
q4_test <- svychisq(~full_professor + sex, d4)
q4_diff_ci <- as.numeric(q4_diff) + c(-1, 1) * 1.96 * as.numeric(SE(q4_diff))

# counts for the text
n_missing <- colSums(is.na(dat[, c("female", "ongoing_grant_2026", "has_phd",
                                   "phd_year", "full_professor")]))

n_phd <- sum(dat$has_phd %in% 1)
n_nophd <- sum(dat$has_phd %in% 0)
n_women <- sum(dat$female %in% 1)
n_men <- sum(dat$female %in% 0)
n_emer <- sum(dat$emeritus)
n_dbl <- sum(dat$double_listed)

# sensitivity: drop emeriti and double-listed faculty from the sample
des2 <- subset(des, !emeritus & !double_listed)
s1 <- svyciprop(~female, subset(des2, !is.na(female)))
s2 <- svyciprop(~ongoing_grant_2026, subset(des2, !is.na(ongoing_grant_2026)))
s3 <- svymean(~years_since_phd, subset(des2, has_phd %in% 1 & !is.na(years_since_phd)))
s4 <- svyciprop(~full_professor, des2)
res <- data.frame(
  Question = c("1. Women", "2. Ongoing grant in 2026",
               "3. Mean years since Ph.D. (Ph.D. holders)",

```

```

      "4a. Full professor, all faculty",
      "4b. Full professor, women",
      "4c. Full professor, men",
      "4d. Difference in full professors (women - men)"),
  Estimate = c(sprintf("%.1f%%", 100 * q1), sprintf("%.1f%%", 100 * q2),
    sprintf("%.1f years", coef(q3)), sprintf("%.1f%%", 100 * q4),
    sprintf("%.1f%%", 100 * q4_by$full_professor[2]),
    sprintf("%.1f%%", 100 * q4_by$full_professor[1]),
    sprintf("%.1f%%", 100 * as.numeric(q4_diff))),
  `95% CI` = c(ci_pct(q1_ci), ci_pct(q2_ci),
    sprintf("%.1f to %.1f years", q3_ci[1], q3_ci[2]),
    ci_pct(q4_ci),
    ci_pct(c(q4_by$ci_l[2], q4_by$ci_u[2])),
    ci_pct(c(q4_by$ci_l[1], q4_by$ci_u[1])),
    sprintf("%.1f%% to %.1f%%", 100 * q4_diff_ci[1], 100 * q4_diff_ci[2])),
  check.names = FALSE)
knitr::kable(res, align = "lrr", caption = "Design-based estimates with 95% confidence intervals")
est <- c(as.numeric(q1), as.numeric(q2), as.numeric(q4),
  q4_by$full_professor[2], q4_by$full_professor[1])
lo <- c(q1_ci[1], q2_ci[1], q4_ci[1], q4_by$ci_l[2], q4_by$ci_l[1])
hi <- c(q1_ci[2], q2_ci[2], q4_ci[2], q4_by$ci_u[2], q4_by$ci_u[1])
labs <- c("Women (Q1)", "Ongoing grant (Q2)", "Full professor, all (Q4)", "Full professor, women")
par(mar = c(4, 9.5, 1, 1))
y <- rev(seq_along(est))
plot(est, y, xlim = c(0, 1), ylim = c(0.5, length(est) + 0.5), pch = 19,
  yaxt = "n", xaxt = "n", xlab = "Estimated proportion", ylab = "")
segments(lo, y, hi, y, lwd = 2)
axis(1, at = seq(0, 1, 0.2), labels = paste0(seq(0, 100, 20), "%"))
axis(2, at = y, labels = labs, las = 1, tick = FALSE)
abline(v = 0.5, lty = 3, col = "grey60")

```